

MD ZEESHAN RAZA

ML & Agentic AI Engineer | Generative AI | LLM Engineer

Bangalore, India | +91 74496 71033 | zshnreza@gmail.com
LinkedIn: linkedin.com/in/zeeshan-ai-engineer | GitHub: github.com/Zshnreza
Portfolio: zeeshan-portfolio-chi.vercel.app | Calendly: calendly.com/zshnreza

PROFESSIONAL SUMMARY

ML and Agentic AI Engineer with 3 years of experience building production-grade Generative AI systems: agentic workflows, RAG pipelines, and multimodal machine learning platforms on AWS Bedrock. Skilled in LLM engineering across the end-to-end ML lifecycle - fine-tuning, LangChain/LangGraph orchestration, contextual chunking, knowledge graphs, evaluation frameworks, and cloud-native deployment with Docker and CI/CD. Enabled 40% faster data-driven decision-making through real-time monitoring and scalable AI pipelines powering 5+ enterprise projects. Open to Machine Learning Engineer, AI Engineer, and Generative AI Engineer roles in Bangalore and globally.

TECHNICAL SKILLS

Generative & Agentic AI: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Agentic AI, Multi-Agent Systems, LangChain, LangGraph, LLM Fine-Tuning, Prompt Engineering, Contextual Chunking, Knowledge Graphs, SPARQL, Vector Databases, Multimodal AI, NLP, NLU, LLM Evaluation (LLM-as-Judge, BLEU, ROUGE, BERTScore)

ML Frameworks: PyTorch, TensorFlow, scikit-learn, Keras, Hugging Face Transformers

Cloud & MLOps: AWS Bedrock, AWS SageMaker, AWS Lambda, AWS S3, Docker, REST APIs (Flask), CI/CD, Git, MLOps, Observability, Agile/Scrum

Data & Analytics: Python, SQL, PostgreSQL, SAS, Alteryx, ETL Pipelines, Feature Engineering, Statistical Analysis

Visualization & BI: Streamlit, Power BI, Tableau, Google Analytics, Mixpanel, Advanced Excel, Retool

WORK EXPERIENCE

ML Data Associate II (ML & Agentic AI Engineer)	Nov 2024 - Mar 2026
Amazon Web Services (AWS) - Bedrock	Bangalore, India
- Engineered scalable ETL pipelines for multimodal LLM training across text, image, audio, and video on AWS Bedrock, supporting 5+ concurrent enterprise projects with improved reliability. - Built Agentic AI systems using LangGraph and LangChain, implementing multi-step reasoning agents, tool-calling pipelines, and RAG workflows with contextual chunking and SPARQL-based knowledge graph integration. - Fine-tuned LLMs (SLMs, LRMs) using PyTorch and TensorFlow, and designed AI quality-assurance frameworks (LLM-as-judge, BLEU, ROUGE, BERTScore) to evaluate accuracy, hallucination, groundedness, and safety at production scale. - Enabled 40% faster data-driven decisions by developing real-time Streamlit dashboards and monitoring tools for pipeline health and model performance, consumed by cross-functional teams. - Productionized enterprise AI workflows by containerizing solutions with Docker, deploying Flask-based REST APIs, and implementing CI/CD, observability, alerting, and runbooks for enterprise-grade reliability.	
Market Research Analyst & Data Analyst	Jul 2023 - Nov 2024
Ayvole	Dehradun, India
- Designed end-to-end data collection and analytics pipelines, aggregating structured datasets from 100+ sources. - Supported large-scale research analyzing 4,000+ startups across 8 sectors for Outlook Business - Start-Up Outperformers 2023. - Conducted market research and gap analysis, translating statistical findings into business recommendations and visual presentations for senior stakeholders. - Drove three new media product launches, directly contributing to a 30% increase in market share across client engagements.	

- Collected, processed, and analyzed operational data while ensuring 100% compliance with industry safety and documentation standards.
- Delivered data-driven recommendations supporting four or more operational departments.

PROJECTS

Agentic AI & RAG Pipeline System (AWS Bedrock)

Designed and deployed a production-grade multi-agent AI system using LangGraph, integrating contextual-chunking RAG pipelines and SPARQL-based knowledge graphs, with full Dockerization, CI/CD pipelines, monitoring, and observability on AWS.

LLM Fine-Tuning & AI Evaluation Framework (AWS Bedrock)

Developed end-to-end LLM fine-tuning pipelines leveraging PyTorch and TensorFlow, implementing offline and online evaluation frameworks (LLM-as-Judge, groundedness, safety, hallucination detection) with automated Python dashboards.

Multimodal AI Data Pipeline (AWS Bedrock)

Architected scalable multimodal data ingestion pipelines across text, image, audio, and video to optimize foundation models, with real-time monitoring dashboards tracking pipeline health and model quality metrics.

Zeeshan Pro Agent (Open Source)

ChatGPT-style personal agent workspace built on LangGraph and LangChain with human-approval gates before external actions; hosted inference (Groq, Hugging Face) and Qdrant Cloud vector storage.

github.com/Zshnreza/zeeshan-pro-agent

RAG Telegram Bot (Open Source, Live)

Fully local RAG chatbot answering questions from custom documents with source attribution: Mistral via Ollama, Sentence Transformers embeddings, FAISS vector search, Telegram Bot API - zero API cost.

t.me/zeeshan_rag_bot

AI Engineer Portfolio (Full-Stack)

Cinematic portfolio built with Next.js 15, React 19, TypeScript, Tailwind CSS v4, and Framer Motion; interactive particle systems, 3D skill sphere, command menu, full SEO and PWA scaffolding, deployed on Vercel.

zeeshan-portfolio-chi.vercel.app

EDUCATION

B.Tech, Electrical & Electronics Engineering

Jun 2020 - May 2023

Swami Vivekananda Institute of Science and Technology, Kolkata - CGPA: 8.4

CERTIFICATIONS

- AWS Educate: Introduction to Cloud 101 - Amazon Web Services (Credly Verified)
- Summarize & Simplify Information with Microsoft 365 Copilot - Microsoft Learn
- Get Started Building with Power BI - Microsoft Learn
- Data Analysis with Python - freeCodeCamp
- Intermediate SQL - DataCamp
- Introduction to SQL - DataCamp
- Introduction to Relational Databases in SQL - DataCamp
- Google Analytics for Beginners - Google Analytics Academy
- AI For India - HCL GUVI
- Google for Education - HCL GUVI
- 2 Years on Kaggle - Kaggle